

CMPT 210: Probability and Computing

Lecture 22

Sharan Vaswani

November 26, 2024

Tail inequalities

- Variance gives us one way to measure how “spread” the distribution is.

Tail inequalities bound the probability that the r.v. takes a value much different from its mean.

Example: Consider a r.v. X that can take on only non-negative values and $\mathbb{E}[X] = 99.99$. Show that $\Pr[X \geq 300] \leq \frac{1}{3}$.

$$\begin{aligned} \text{Proof: } \mathbb{E}[X] &= \sum_{x \in \text{Range}(X)} x \Pr[X = x] = \sum_{x|x \geq 300} x \Pr[X = x] + \sum_{x|0 \leq x < 300} x \Pr[X = x] \\ &\geq \sum_{x|x \geq 300} (300) \Pr[X = x] + \sum_{x|0 \leq x < 300} x \Pr[X = x] \\ &= (300) \Pr[X \geq 300] + \sum_{x|0 \leq x < 300} x \Pr[X = x] \end{aligned}$$

If $\Pr[X \geq 300] > \frac{1}{3}$, then, $\mathbb{E}[X] > (300) \frac{1}{3} + \sum_{x|0 \leq x < 300} x \Pr[X = x] > 100$ (since the second term is always non-negative). Hence, if $\Pr[X \geq 300] > \frac{1}{3}$, $\mathbb{E}[X] > 100$ which is a contradiction since $\mathbb{E}[X] = 99.99$.

Markov's Theorem

Markov's theorem formalizes the intuition on the last slide of the previous class, and can be stated as follows.

Markov's Theorem: If X is a non-negative random variable, then for all $x > 0$,

$$\Pr[X \geq x] \leq \frac{\mathbb{E}[X]}{x}.$$

Proof: Define $\mathcal{I}\{X \geq x\}$ to be the indicator r.v. for the event $[X \geq x]$. Then for all values of X , $x\mathcal{I}\{X \geq x\} \leq X$.

$$\begin{aligned}\mathbb{E}[x\mathcal{I}\{X \geq x\}] &\leq \mathbb{E}[X] \implies x\mathbb{E}[\mathcal{I}\{X \geq x\}] \leq \mathbb{E}[X] \implies x\Pr[X \geq x] \leq \mathbb{E}[X] \\ &\implies \Pr[X \geq x] \leq \frac{\mathbb{E}[X]}{x}.\end{aligned}$$

Since the above theorem holds for all $x > 0$, we can set $x = c\mathbb{E}[X]$ for $c \geq 1$. In this case, $\Pr[X \geq c\mathbb{E}[X]] \leq \frac{1}{c}$. Hence, the probability that X is “far” from the mean in terms of the multiplicative factor c is upper-bounded by $\frac{1}{c}$.

Markov's Theorem – Example

Q: If X is a non-negative r.v. such that $\mathbb{E}[X] = 150$, compute the probability that X is at least 200.

Q: If we are provided additional information that X can not take values less than 100 and $\mathbb{E}[X] = 150$, compute the probability that X is at least 200.

Define $Y := X - 100$. $\mathbb{E}[Y] = \mathbb{E}[X] - 100 = 50$ and Y is non-negative.

$$\Pr[X \geq 200] = \Pr[Y + 100 \geq 200] = \Pr[Y \geq 100] \leq \frac{\mathbb{E}[Y]}{100} = \frac{50}{100} = \frac{1}{2}$$

Hence, if we have additional information (in the form of a lower-bound that a r.v. can not be smaller than some constant $b > 0$), we can use Markov's Theorem on the shifted r.v. (Y in our example) and obtain a tighter bound on the probability of deviation.

Chebyshev's Theorem

Chebyshev's Theorem: For a r.v. X and any constant $y > 0$,

$$\Pr[|X - \mathbb{E}[X]| \geq y] \leq \frac{\text{Var}[X]}{y^2}.$$

Proof: Use Markov's Theorem with some cleverly chosen function of X . Formally, for some function f such that $Y := f(X)$ is non-negative. Using Markov's Theorem for Y ,

$$\Pr[f(X) \geq x] \leq \frac{\mathbb{E}[f(X)]}{x}$$

Choosing $f(X) = |X - \mathbb{E}[X]|^2$ and $x = y^2$ implies that $f(X)$ is non-negative and $x > 0$. Using Markov's Theorem,

$$\Pr[|X - \mathbb{E}[X]|^2 \geq y^2] \leq \frac{\mathbb{E}[|X - \mathbb{E}[X]|^2]}{y^2}$$

Note that $\Pr[|X - \mathbb{E}[X]|^2 \geq y^2] = \Pr[|X - \mathbb{E}[X]| \geq y]$, and hence,

$$\Pr[|X - \mathbb{E}[X]| \geq y] \leq \frac{\mathbb{E}[|X - \mathbb{E}[X]|^2]}{y^2} = \frac{\text{Var}[X]}{y^2}$$

Chebyshev's Theorem

- Chebyshev's Theorem bounds the probability that the random variable X is “far” away from the mean $\mathbb{E}[X]$ by an additive factor of x .
- If we set $x = c\sigma_X$ where σ_X is the standard deviation of X , then by Chebyshev's Theorem,

$$\Pr[(X \geq \mathbb{E}[X] + c\sigma_X) \cup (X \leq \mathbb{E}[X] - c\sigma_X)] = \Pr[|X - \mathbb{E}[X]| \geq c\sigma_X] \leq \frac{\text{Var}[X]}{c^2\sigma_X^2} = \frac{1}{c^2}$$

$$\Pr[\mathbb{E}[X] - c\sigma_X < X < \mathbb{E}[X] + c\sigma_X] = \Pr[|X - \mathbb{E}[X]| \leq c\sigma_X]$$

$$\implies \Pr[\mathbb{E}[X] - c\sigma_X < X < \mathbb{E}[X] + c\sigma_X] = 1 - \Pr[|X - \mathbb{E}[X]| \geq c\sigma_X] \geq 1 - \frac{1}{c^2}.$$

Hence, Chebyshev's Theorem can be used to bound the probability that X is “concentrated” near its mean.

- Unlike Markov's Theorem, Chebyshev's Theorem does not require the r.v. to be non-negative, but requires knowledge of the variance.

Chebyshev's Theorem - Example

Q: If X is a non-negative r.v. such that $\mathbb{E}[X] = 100$ and $\sigma_X = 15$, compute the probability that X is at least 300.

If we use Markov's Theorem, $\Pr[X \geq 300] \leq \frac{\mathbb{E}[X]}{300} = \frac{1}{3}$.


Note that $\Pr[|X - 100| \geq 200] = \Pr[X \leq -100 \cup X \geq 300] = \Pr[X \geq 300]$. Using Chebyshev's Theorem,

$$\Pr[X \geq 300] = \Pr[|X - 100| \geq 200] \leq \frac{\text{Var}[X]}{(200)^2} = \frac{15^2}{200^2} \approx \frac{1}{178}.$$

Hence, by exploiting the knowledge of the variance and using Chebyshev's inequality, we can obtain a tighter bound.

Chebyshev's Theorem - Example

Q: Consider a r.v. $X \sim \text{Bin}(20, 0.75)$. Plot the PDF_X , compute its mean and standard deviation and bound $\Pr[10 < X < 20]$.



$\text{Range}(X) = \{0, 1, \dots, 20\}$ and for $k \in \text{Range}(X)$,
 $f(k) = \binom{n}{k} p^k (1-p)^{n-k}$.

$$\mathbb{E}[X] = np = (20)(0.75) = 15$$

$$\text{Var}[X] = np(1-p) = 20(0.75)(0.25) = 3.75 \text{ and hence } \sigma_X = \sqrt{3.75} \approx 1.94.$$

$$\begin{aligned} \Pr[10 < X < 20] &= 1 - \Pr[X \leq 10 \cup X \geq 20] \\ &= 1 - \Pr[|X - 15| \geq 5] \\ &= 1 - \Pr[|X - \mathbb{E}[X]| \geq 5] \\ &\geq 1 - \frac{\text{Var}[X]}{(5)^2} = 1 - \frac{3.75}{25} = 0.85. \end{aligned}$$

Hence, the “probability mass” of X is “concentrated” around its mean.

Voter Poll

Q: Suppose there is an election between two candidates Donald Trump and Kamala Harris, and we are hired by candidate Harris' election campaign to estimate his chances of winning the election. In particular, we want to estimate p , the fraction of voters favoring Harris before the election. We conduct a voter poll – selecting (typically calling) people uniformly at random (with replacement so that we can choose a person twice) and try to estimate p . What is the number of people we should poll to estimate p reasonably accurately and with reasonably high probability?

Define X_i to be the indicator r.v. equal to 1 iff person i that we called favors Harris.

Assumption (1): The X_i r.v.'s are mutually independent since the people we poll are chosen randomly and we assume that their opinions do not affect each other.

Assumption (2): The people we call are identically distributed i.e. $X_i = 1$ with probability p .

Suppose we poll n people and define $S_n := \sum_{i=1}^n X_i$ as the r.v. equal to the total number of people (amongst the ones we polled) that prefer Harris. $\frac{S_n}{n}$ is the *statistical estimate* of p .

Q: What is the distribution of S_n ?

Goal: We want to find for what n is our estimate for p accurate up to an error $\epsilon > 0$ and with probability $1 - \delta$ (for $\delta \in (0, 1)$). Formally, we want to find an n such that

$$\Pr \left[\left| \frac{S_n}{n} - p \right| < \epsilon \right] \geq 1 - \delta.$$

Since $S_n \sim \text{Bin}(n, p)$, $\mathbb{E}[S_n] = np$ and hence, $\mathbb{E} \left[\frac{S_n}{n} \right] = p$, meaning that our estimate is *unbiased* – in expectation, the estimate is equal to p . Hence, the above statement is equivalent to,

$$\Pr \left[\left| \frac{S_n}{n} - \mathbb{E} \left[\frac{S_n}{n} \right] \right| < \epsilon \right] \geq 1 - \delta$$

Hence, we can use Chebyshev's Theorem for the r.v. $\frac{S_n}{n}$ with $x = \epsilon$ to bound the LHS

$$\Pr \left[\left| \frac{S_n}{n} - \mathbb{E} \left[\frac{S_n}{n} \right] \right| < \epsilon \right] = 1 - \Pr \left[\left| \frac{S_n}{n} - \mathbb{E} \left[\frac{S_n}{n} \right] \right| \geq \epsilon \right] \geq 1 - \frac{\text{Var}[S_n/n]}{\epsilon^2}.$$

In order to achieve our goal, it is sufficient to find an n such that,

$$1 - \frac{\text{Var}[S_n/n]}{\epsilon^2} \geq 1 - \delta \implies \frac{\text{Var}[S_n/n]}{\epsilon^2} \leq \delta$$

Voter Poll

Let us calculate the $\text{Var}[S_n/n]$.

$$\begin{aligned}\text{Var}[S_n/n] &= \frac{1}{n^2} \text{Var}[S_n] && \text{(Using the property of variance)} \\ &= \frac{1}{n^2} n p (1-p) = \frac{p(1-p)}{n} && \text{(Using the variance of the Binomial distribution)}\end{aligned}$$

Hence, we want to find n s.t.

$$\frac{p(1-p)}{n\epsilon^2} \leq \delta \implies n \geq \frac{p(1-p)}{\epsilon^2 \delta}$$

But we do not know p ! If $n \geq \max_p \frac{p(1-p)}{\epsilon^2 \delta}$, then for any p , $n \geq \frac{p(1-p)}{\epsilon^2 \delta}$. So the problem is to compute $\max_p \frac{p(1-p)}{\epsilon^2 \delta}$. This is a concave function and is maximized at $p = 1/2$.

Hence, $n \geq \frac{1}{4\epsilon^2 \delta}$ is sufficient to ensure that $\Pr[|\frac{S_n}{n} - p| < \epsilon] \geq 1 - \delta$ meaning that we have estimated p upto an error ϵ and this bound is true with high probability equal to $1 - \delta$.

For example, if $\epsilon = 0.01$ and $\delta = 0.01$ meaning that we want the bound to hold 99% of the time, then, we require $n \geq 250000$.

Questions?