

CMPT 409/981: Optimization for Machine Learning

Lecture 19

Sharan Vaswani

November 21, 2022

Min-Max Optimization

Today we will focus on problems of the form

$$\min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v).$$

Application: Two player zero-sum matrix games of the form,

$$\min_{w \in \Delta_A} \max_{v \in \Delta_B} w^\top M v,$$

where A is the set of strategies available to player 1. $\Delta_A = \{w \in [0, 1]^{|A|} \mid \sum_i w_i = 1\}$ is the distribution over these available strategies and $w \in \Delta_A$ is a possible **mixed strategy**.

The matrix $M \in \mathbb{R}^{|A| \times |B|}$ is the **payoff matrix** for player 1 i.e. if player 1 plays strategy i and player 2 plays strategy j , then player 1 is penalized $M_{i,j}$ whereas player 2 is penalized $-M_{i,j}$. Both players are trying to minimize their respective penalties.

Since (penalty for player 1) = -(penalty for player 2), this is a zero-sum game. Classic example: Rock-Paper-Scissors

Min-Max Optimization

Application: Generative Adversarial Networks

$$\min_{\theta} \max_{\phi} [\mathbb{E}_{x \sim p_{\text{real}}} [\log D_{\phi}(x)] + \mathbb{E}_{z \sim N(0, I_d)} [\log (1 - D_{\phi}(G_{\theta}(z)))]],$$

where $G_{\theta}(z)$ is the generator parameterized by θ that attempts to generate realistic images from random noise z . $D_{\phi}(x)$ is the discriminator parameterized by ϕ that attempts to discriminate between the real (from p_{real}) and generated (from $G_{\theta}(z)$) images.

Application: Distributionally Robust Optimization

$$\min_{\theta} \max_{P \in \mathcal{P}} \mathbb{E}_{\zeta \in P} \ell(\theta, \zeta)$$

where $\mathcal{P} := \{P | d(P, \hat{P}) \leq \rho\}$ is the family of distributions that are “close” (measured by ρ) to the empirical distribution \hat{P} according to a distance metric d (Total variation, KL divergence).

We require that the model (parameterized by θ) is robust to distributions close to the empirical distribution from which can obtain samples.

Min-Max Optimization

Let us abstract out these problems and consider the following objective,

$$\min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$$

where $\mathcal{W} \subseteq \mathbb{R}^{d_w}$ and $\mathcal{V} \subseteq \mathbb{R}^{d_v}$ are convex sets.

Claim: In general, $\max_{v \in \mathcal{V}} \min_{w \in \mathcal{W}} f(w, v) \leq \min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$

Proof: Define $v^* := \arg \max_{v \in \mathcal{V}} \min_{w \in \mathcal{W}} f(w, v)$ and $w^* := \arg \min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$.

$$\max_{v \in \mathcal{V}} \min_{w \in \mathcal{W}} f(w, v) = \min_{w \in \mathcal{W}} f(w, v^*) \leq f(w^*, v^*) \leq \max_{v \in \mathcal{V}} f(w^*, v) = \min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$$

Game theoretic interpretation: RHS corresponds to w -player playing first and the v -player reacting, while the LHS corresponds to the v -player playing first and the w -player reacting. Since the v -player aims to maximize f , playing second might be beneficial since they can adapt to the w -player's strategy. Hence, the $\text{RHS} \geq \text{LHS}$.

Min-Max Optimization

Convex-Concave Games: $f : \mathcal{W} \times \mathcal{V} \rightarrow \mathbb{R}$ is convex-concave iff $f(\cdot, v)$ is a convex function for any $v \in \mathcal{V}$, $f(w, \cdot)$ is a concave function for any $w \in \mathcal{W}$ and \mathcal{W}, \mathcal{V} are convex sets.

Sion's Minimax Theorem: If \mathcal{W} and \mathcal{V} are compact, convex sets, and f is a convex-concave function, then $\max_{v \in \mathcal{V}} \min_{w \in \mathcal{W}} f(w, v) = \min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$.

Example: $f(w, v) = \min_{w \in \Delta_A} \max_{v \in \Delta_B} w^T M v$ is convex-concave and the simplex Δ is a convex set. Hence it is a convex-concave game.

Recall that $v^* := \arg \max_{v \in \mathcal{V}} \min_{w \in \mathcal{W}} f(w, v)$ and $w^* := \arg \min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$ and

$$\max_{v \in \mathcal{V}} \min_{w \in \mathcal{W}} f(w, v) = \min_{w \in \mathcal{W}} f(w, v^*) \leq f(w^*, v^*) \leq \max_{v \in \mathcal{V}} f(w^*, v) = \min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$$

If f convex-concave and \mathcal{W} and \mathcal{V} are convex sets, then,

$$\max_{v \in \mathcal{V}} \min_{w \in \mathcal{W}} f(w, v) = \min_{w \in \mathcal{W}} f(w, v^*) = f(w^*, v^*) = \max_{v \in \mathcal{V}} f(w^*, v) = \min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v).$$

Hence, (w^*, v^*) is a solution to the game iff for all $w \in \mathcal{W}$, $v \in \mathcal{V}$,

$$f(w^*, v) \leq f(w^*, v^*) \leq f(w, v^*).$$

Min-Max Optimization

Recall that for convex-concave games, (w^*, v^*) is a solution iff for all $w \in \mathcal{W}$, $v \in \mathcal{V}$, $f(w^*, v) \leq f(w^*, v^*) \leq f(w, v^*)$.

Game theoretic interpretation: From the perspective of a game between the w -player and the v -player, since $f(w^*, v^*) = \min_{w \in \mathcal{W}} f(w, v^*)$, if the v -player is playing v^* , it is optimal for the w -player to play w^* . Similarly, if the w -player is playing w^* , it is optimal for the v -player to play v^* . Hence, (w^*, v^*) is the **Nash equilibrium** since neither player has an incentive to move away from their strategy.

For convex-concave games, the Nash equilibrium is guaranteed to exist, but need not be unique.

Duality Gap: Way to characterize the sub-optimality of the point (\hat{w}, \hat{v}) :

$$\text{Duality Gap}((\hat{w}, \hat{v})) := \max_{v \in \mathcal{V}} f(\hat{w}, v) - \min_{w \in \mathcal{W}} f(w, \hat{v}).$$

If (\hat{w}, \hat{v}) is a Nash equilibrium, then $\max_{v \in \mathcal{V}} f(\hat{w}, v) = f(\hat{w}, \hat{v}) = \min_{w \in \mathcal{W}} f(w, \hat{v})$ and hence the duality gap is 0. Point (\hat{w}, \hat{v}) is an ϵ -Nash equilibrium, if the $\text{Duality Gap}((\hat{w}, \hat{v})) \leq \epsilon$.

Questions?

Gradient Descent Ascent

Gradient Descent Ascent: At iteration k , for a step-size η , (simultaneous) projected Gradient Descent Ascent (GDA) has the following update:

$$w_{k+1} = \Pi_{\mathcal{W}}[w_k - \eta_k \nabla_w f(w_k, v_k)] \quad ; \quad v_{k+1} = \Pi_{\mathcal{V}}[v_k + \eta_k \nabla_v f(w_k, v_k)] ,$$

where $\Pi_{\mathcal{W}}$ and $\Pi_{\mathcal{V}}$ are Euclidean projections onto \mathcal{W} and \mathcal{V} respectively. It is possible to use different step-sizes for the w and v variables.

G-Lipschitz functions: Define $z = \begin{bmatrix} w \\ v \end{bmatrix}$. The function $f : \mathcal{W} \times \mathcal{V} \rightarrow \mathbb{R}$ is G -Lipschitz iff,

$$|f(z_1) - f(z_2)| \leq G \|z_1 - z_2\|$$

Similar to convex minimization, this implies bounded gradients, i.e. for all $w \in \mathcal{W}$, $v \in \mathcal{V}$,

$$\|\nabla_w f(w, v)\| \leq G \quad ; \quad \|\nabla_v f(w, v)\| \leq G$$

We will also assume that sets \mathcal{W} and \mathcal{V} have diameter D i.e. for all $w_1, w_2 \in \mathcal{W}$, $\|w_1 - w_2\|^2 \leq D^2$. Similarly, for all $v_1, v_2 \in \mathcal{V}$, $\|v_1 - v_2\|^2 \leq D^2$.

Gradient Descent Ascent for Lipschitz, convex-concave games

Claim: For G -Lipschitz convex-concave games where \mathcal{W} and \mathcal{V} have diameter D , projected GDA with $\eta_k = \frac{D}{\sqrt{2G\sqrt{k}}}$ results in the following bound for $\bar{w}_T := \sum_{k=1}^T w_k / T$ and $\bar{v}_T := \sum_{k=1}^T v_k / T$

$$\text{Duality Gap}((\bar{w}_T, \bar{v}_T)) \leq \frac{4DG}{\sqrt{T}}$$

Proof: For a fixed point $\tilde{w} \in \mathcal{W}$, using the projected gradient descent update for w ,

$$\begin{aligned} \|w_{k+1} - \tilde{w}\|^2 &= \|\Pi_{\mathcal{W}}[w_k - \eta \nabla_w f(w_k, v_k)] - \Pi_{\mathcal{W}}[\tilde{w}]\|^2 && \text{(Since } \tilde{w} \in \mathcal{W}\text{)} \\ &\leq \|w_k - \eta \nabla_w f(w_k, v_k) - \tilde{w}\|^2 \end{aligned}$$

(since projections are non-expansive)

$$\begin{aligned} &= \|w_k - \tilde{w}\|^2 - 2\eta_k \langle \nabla_w f(w_k, v_k), w_k - \tilde{w} \rangle + \eta_k^2 \|\nabla_w f(w_k, v_k)\|^2 \\ &\leq \|w_k - \tilde{w}\|^2 - 2\eta_k [f(w_k, v_k) - f(\tilde{w}, v_k)] + \eta_k^2 G^2 \end{aligned}$$

(Since $f(\cdot, v_k)$ is convex and f is G -Lipschitz)

$$\implies [f(w_k, v_k) - f(\tilde{w}, v_k)] \leq \frac{\|w_k - \tilde{w}\|^2 - \|w_{k+1} - \tilde{w}\|^2}{2\eta_k} + \frac{\eta_k}{2} G^2 \quad (1)$$

Gradient Descent Ascent for Lipschitz, convex-concave games

Similarly, using the projected gradient ascent update w.r.t $\tilde{v} \in \mathcal{V}$,

$$\begin{aligned}\|v_{k+1} - \tilde{v}\|^2 &\leq \|v_k - \tilde{v}\|^2 + 2\eta_k \langle \nabla_v f(w_k, v_k), v_k - \tilde{v} \rangle + \eta_k^2 \|\nabla_v f(w_k, v_k)\|^2 \\ &\leq \|v_k - \tilde{v}\|^2 + 2\eta_k [f(w_k, v_k) - f(w_k, \tilde{v})] + \eta_k^2 G^2 \\ &\quad \text{(Since } f(w_k, \cdot) \text{ is concave and } f \text{ is } G\text{-Lipschitz)}\end{aligned}$$

$$\Rightarrow [f(w_k, \tilde{v}) - f(w_k, v_k)] \leq \frac{\|v_k - \tilde{v}\|^2 - \|v_{k+1} - \tilde{v}\|^2}{2\eta_k} + \frac{\eta_k}{2} G^2 \quad (2)$$

Adding Eq. (1) and Eq. (2),

$$\begin{aligned}f(w_k, \tilde{v}) - f(\tilde{w}, v_k) &\leq \frac{\|w_k - w\|^2 - \|w_{k+1} - w\|^2}{2\eta_k} + \frac{\|v_k - v\|^2 - \|v_{k+1} - v\|^2}{2\eta_k} + \eta_k G^2 \\ \sum_{k=1}^T [f(w_k, \tilde{v}) - f(\tilde{w}, v_k)] &\leq \sum_{k=1}^T \left[\frac{\|w_k - \tilde{w}\|^2 - \|w_{k+1} - \tilde{w}\|^2}{2\eta_k} \right] + \sum_{k=1}^T \left[\frac{\|v_k - \tilde{v}\|^2 - \|v_{k+1} - \tilde{v}\|^2}{2\eta_k} \right] \\ &\quad + G^2 \sum_{k=1}^T \eta_k\end{aligned}$$

Gradient Descent Ascent for Lipschitz, convex-concave games

Simplifying the first term in the equation from the previous slide,

$$\begin{aligned}\sum_{k=1}^T \left[\frac{\|w_k - \tilde{w}\|^2 - \|w_{k+1} - \tilde{w}\|^2}{2\eta_k} \right] &= \sum_{k=2}^{T-1} \|w_k - \tilde{w}\|^2 \left[\frac{1}{\eta_{k-1}} - \frac{1}{\eta_k} \right] + \frac{\|w_1 - w^*\|^2}{2\eta_1} - \frac{\|w_T - w^*\|^2}{2\eta_T} \\ &\leq \frac{D^2}{2\eta_T}\end{aligned}$$

Bounding the second term in a similar manner and putting everything together,

$$\begin{aligned}\sum_{k=1}^T [f(w_k, \tilde{v}) - f(\tilde{w}, v_k)] &\leq \frac{D^2}{\eta_T} + G^2 \sum_{k=1}^T \eta_k = \frac{D^2 \sqrt{T}}{\eta} + G^2 \eta \sum_{k=1}^T \frac{1}{\sqrt{k}} \\ &\hspace{25em} (\eta_k = \eta/\sqrt{k}) \\ &\leq \frac{D^2 \sqrt{T}}{\eta} + 2G^2 \eta \sqrt{T} = 2\sqrt{2} DG \sqrt{T} \hspace{2em} (\eta = \frac{D}{\sqrt{2G}}) \\ \implies \frac{1}{T} \left[\sum_{k=1}^T [f(w_k, \tilde{v}) - f(\tilde{w}, v_k)] \right] &\leq \frac{4 DG}{\sqrt{T}}\end{aligned}$$

Gradient Descent Ascent for Lipschitz, convex-concave games

Recall that $\frac{1}{T} \left[\sum_{k=1}^T [f(w_k, \tilde{v}) - f(\tilde{w}, v_k)] \right] \leq \frac{4DG}{\sqrt{T}}$. Since $f(\cdot, \tilde{y})$ and $-f(\tilde{w}, \cdot)$ are convex, using Jensen's inequality and by definition of \bar{w}_T and \bar{v}_T ,

$$f(\bar{w}_T, \tilde{v}) - f(\tilde{w}, \bar{v}_T) \leq \frac{4DG}{\sqrt{T}}$$

Since the above statement is true for all $\tilde{v} \in \mathcal{V}$ and $\tilde{w} \in \mathcal{W}$, taking the maximum over $\tilde{v} \in \mathcal{V}$ and the minimum over $\tilde{w} \in \mathcal{W}$,

$$\max_{v \in \mathcal{V}} f(\bar{w}_T, v) - \min_{w \in \mathcal{W}} f(w, \bar{v}_T) \leq \frac{4DG}{\sqrt{T}} \implies \text{Duality Gap}((\bar{w}_T, \bar{v}_T)) \leq \frac{4DG}{\sqrt{T}}$$

Recall that GD attains an $O(1/\sqrt{T})$ rate when minimizing convex, Lipschitz functions, and hence GDA has a similar behaviour when solving convex-concave Lipschitz games.

Questions?

Gradient Descent Ascent for smooth, convex-concave games

Similar to convex minimization, $f : \mathcal{W} \times \mathcal{V} \rightarrow \mathbb{R}$ is L -smooth iff

$$\|\nabla_w f(w_1, v_1) - \nabla_w f(w_2, v_2)\| \leq L \|z_1 - z_2\| \quad ; \quad \|\nabla_v f(w_1, v_1) - \nabla_v f(w_2, v_2)\| \leq L \|z_1 - z_2\| ,$$

where $z_1 = \begin{bmatrix} w_1 \\ v_1 \end{bmatrix}$ and $z_2 = \begin{bmatrix} w_2 \\ v_2 \end{bmatrix}$.

Example: $f(w, v) = wv$ is 1-smooth since $\nabla_w f(w, v) = v$ and

$|v_1 - v_2| \leq |v_1 - v_2| + |w_1 - w_2|$. A similar reasoning works for $\nabla_v f(w, v)$. Moreover, since $f(\cdot, v)$ is linear w.r.t w , it is convex. By symmetry, $f(w, \cdot)$ is linear in v and hence concave.

If $\mathcal{W} = \mathbb{R}$ and $\mathcal{V} = \mathbb{R}$, $\min_{w \in \mathbb{R}} \max_{v \in \mathbb{R}} wv$ is a smooth, convex-concave game whose unique solution is at $(0, 0)$ since $f(0, 0) \leq f(w, 0)$ for all w and $f(0, 0) \geq f(0, v)$ for all v .

Game theoretically, if the v -player deviates from 0 such that $v = \epsilon$, the w -player can choose $-\infty$ to make the objective small. Similarly, if the w -player deviates from 0 such $w = \epsilon$, then the v -player can choose $+\infty$ to make the objective large. Hence, neither player has an incentive to deviate from $(0, 0)$ which corresponds to the Nash equilibrium.

Gradient Descent Ascent for smooth, convex-concave games

Let us consider running GDA for $\min_{w \in \mathbb{R}} \max_{v \in \mathbb{R}} wv$. The update can be given as:

$$w_{k+1} = w_k - \eta_k \nabla_w f(w_k, v_k) = w_k - \eta_k v_k \quad ; \quad v_{k+1} = v_k + \eta_k \nabla_v f(w_k, v_k) = v_k + \eta_k w_k$$

Calculating the distance from the solution $(0, 0)$ after one iteration,

$$(w_{k+1} - 0)^2 + (v_{k+1} - 0)^2 = (w_k - \eta_k v_k)^2 + (v_k + \eta_k w_k)^2 = (1 + \eta_k^2) (w_k^2 + v_k^2)$$

Hence, for any η_k , the last iterate of GDA will move away from the solution, diverging in the unconstrained setting or hitting the boundary in the constrained setting.

Compare this to GD for smooth, convex minimization where the sub-optimality corresponding to the last iterate decreases at an $O(1/T)$ rate (Lecture 4).

Gradient Descent Ascent for smooth, convex-concave games

Consider a smooth, convex-concave game $\min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$ where the convex sets \mathcal{W} and \mathcal{V} have diameter D . In this case, the duality gap for the *average iterate* of GDA will decrease at a slower $O(1/\sqrt{T})$ rate.

Claim: An L -smooth game $\min_{w \in \mathcal{W}} \max_{v \in \mathcal{V}} f(w, v)$ where \mathcal{W} and \mathcal{V} have diameter D is $\sqrt{2}D$ L -Lipschitz.

Proof: By the definition of L -smoothness,

$$\|\nabla_w f(w_1, v_1) - \nabla_w f(w_2, v_2)\| \leq L \|z_1 - z_2\| \leq L \sqrt{\|w_1 - w_2\|^2 + \|v_1 - v_2\|^2} \leq \sqrt{2}DL.$$

Claim: For L -smooth, convex-concave games, GDA with $\eta_k = \frac{1}{2L\sqrt{k}}$ results in the following bound for $\bar{w}_T := \sum_{k=1}^T w_k / T$ and $\bar{v}_T := \sum_{k=1}^T v_k / T$

$$\text{Duality Gap}((\bar{w}_T, \bar{v}_T)) \leq \frac{4 D^2 L}{\sqrt{T}}$$

Proof: Using the result for convex-concave G -Lipschitz games with $G = \sqrt{2}D L$.

Questions?

Strongly-convex strongly-concave games

Strongly-convex strongly-concave games: $f : \mathcal{W} \times \mathcal{V} \rightarrow \mathbb{R}$ is strongly-convex strongly-concave iff $f(\cdot, v)$ is a strongly-convex function for any $v \in \mathcal{V}$, $f(w, \cdot)$ is a strongly-concave function for any $w \in \mathcal{W}$ and the sets \mathcal{W}, \mathcal{V} are convex sets, i.e. for all $w, w_1, w_2 \in \mathcal{W}$ and $v, v_1, v_2 \in \mathcal{V}$,

$$\begin{aligned} f(w_2, v) &\geq f(w_1, v) + \langle \nabla_w f(w_1, v), w_2 - w_1 \rangle + \frac{\mu_w}{2} \|w_1 - w_2\|^2 \\ -f(w, v_2) &\geq -f(w, v_1) + \langle -\nabla_v f(w, v_1), v_2 - v_1 \rangle + \frac{\mu_v}{2} \|v_1 - v_2\|^2 \end{aligned}$$

If $\mathcal{W} = \mathbb{R}^d$ and $\mathcal{V} = \mathbb{R}^d$ since $w^* := \arg \min_w f(w, v^*)$, $\nabla_w f(w^*, v^*) = 0$. By the strong-convexity of $f(\cdot, v)$ with $v = v^*$, $w_1 = w^*$, $w_2 = w$, $f(w, v^*) > f(w^*, v^*) + \langle \nabla_w f(w^*, v^*), w - w^* \rangle$. Hence, $f(w^*, v^*) < f(w, v^*)$ for all w .

Similarly, $v^* := \arg \max_v f(w^*, v)$, $\nabla_v f(w^*, v^*) = 0$. By the strong-concavity of $f(w, \cdot)$ with $w = w^*$, $-f(w^*, v) > -f(w^*, v^*)$. Hence, $f(w^*, v^*) > f(w^*, v)$ for all v .

Hence, for unconstrained strongly-convex strongly-concave games, (w^*, v^*) is the unique Nash equilibrium and $\nabla_w f(w^*, v^*) = \nabla_v f(w^*, v^*) = 0$.

Gradient Descent Ascent for smooth, strongly-convex strongly-concave games

Let us define an **operator** $F : \mathbb{R}^{d_w+d_v} \rightarrow \mathbb{R}^{d_w+d_v}$ such that the GDA update for unconstrained games can be written as:

$$z_{k+1} = z_k - \eta_k F(z_k) \quad \text{where } z = \begin{bmatrix} w \\ v \end{bmatrix} \quad \text{and } F(z) = \begin{bmatrix} \nabla_w f(w, v) \\ -\nabla_v f(w, v) \end{bmatrix}$$

Recall that in the unconstrained setting, when $\mathcal{W} = \mathbb{R}^{d_w}$ and $\mathcal{V} = \mathbb{R}^{d_v}$, $F(z^*) = 0$.

Claim: If f is L -smooth, then F is $2L$ -Lipschitz i.e. $\|F(z_1) - F(z_2)\| \leq 2L \|z_1 - z_2\|$.

Proof:

$$\begin{aligned} \|F(z_1) - F(z_2)\| &= \left\| \begin{bmatrix} \nabla_w f(w_1, v_1) - \nabla_w f(w_2, v_2) \\ \nabla_v f(w_2, v_2) - \nabla_v f(w_1, v_1) \end{bmatrix} \right\| \\ &\leq \|\nabla_w f(w_1, v_1) - \nabla_w f(w_2, v_2)\| + \|\nabla_v f(w_1, v_1) - \nabla_v f(w_2, v_2)\| \\ &\leq L \|z_1 - z_2\| + L \|z_1 - z_2\| \quad (\text{By definition of } L\text{-smoothness}) \\ \|F(z_1) - F(z_2)\| &\leq 2L \|z_1 - z_2\| \end{aligned}$$

Gradient Descent Ascent for smooth, strongly-convex strongly-concave games

Claim: If f is strongly-convex strongly-concave with $\mu_w = \mu_v = \mu$, then F is μ strongly-monotone i.e. $\langle F(z_1) - F(z_2), z_1 - z_2 \rangle \geq \mu \|z_1 - z_2\|^2$.

Proof: By strong-convexity of $f(\cdot, v)$,

$$f(w_2, v_1) \geq f(w_1, v_1) + \langle \nabla_w f(w_1, v_1), w_2 - w_1 \rangle + \frac{\mu}{2} \|w_1 - w_2\|^2 \quad (\text{With } v = v_1)$$

$$f(w_1, v_2) \geq f(w_2, v_2) + \langle \nabla_w f(w_2, v_2), w_1 - w_2 \rangle + \frac{\mu}{2} \|w_1 - w_2\|^2 \quad (\text{With } v = v_2)$$

By strong-concavity of $f(w, \cdot)$,

$$-f(w_1, v_2) \geq -f(w_1, v_1) + \langle -\nabla_v f(w_1, v_1), v_2 - v_1 \rangle + \frac{\mu}{2} \|v_1 - v_2\|^2 \quad (\text{With } w = w_1)$$

$$-f(w_2, v_1) \geq -f(w_2, v_2) + \langle -\nabla_v f(w_2, v_2), v_1 - v_2 \rangle + \frac{\mu}{2} \|v_1 - v_2\|^2 \quad (\text{With } w = w_2)$$

Adding all the 4 equations,

$$\begin{aligned} & \langle \nabla_w f(w_1, v_1) - \nabla_w f(w_2, v_2), w_1 - w_2 \rangle + \langle \nabla_v f(w_2, v_2) - \nabla_v f(w_1, v_1), v_1 - v_2 \rangle \\ & \geq \mu [\|w_1 - w_2\|^2 + \|v_1 - v_2\|^2] = \mu \|z_1 - z_2\|^2 \end{aligned}$$

Gradient Descent Ascent for smooth, strongly-convex strongly-concave games

Rewriting

$$\langle \nabla_w f(w_1, v_1) - \nabla_w f(w_2, v_2), w_1 - w_2 \rangle + \langle \nabla_v f(w_2, v_2) - \nabla_v f(w_1, v_1), v_1 - v_2 \rangle \geq \mu \|z_1 - z_2\|^2,$$

$$\left\langle \begin{bmatrix} \nabla_w f(w_1, v_1) - \nabla_w f(w_2, v_2) \\ -\nabla_v f(w_1, v_1) + \nabla_v f(w_2, v_2) \end{bmatrix}, \begin{bmatrix} w_1 - w_2 \\ v_1 - v_2 \end{bmatrix} \right\rangle \geq \mu \|z_1 - z_2\|^2$$

$$\left\langle \begin{bmatrix} \nabla_w f(w_1, v_1) \\ -\nabla_v f(w_1, v_1) \end{bmatrix} - \begin{bmatrix} \nabla_w f(w_2, v_2) \\ -\nabla_v f(w_2, v_2) \end{bmatrix}, \begin{bmatrix} w_1 \\ v_1 \end{bmatrix} - \begin{bmatrix} w_2 \\ v_2 \end{bmatrix} \right\rangle \geq \mu \|z_1 - z_2\|^2$$

$$\implies \langle F(z_1) - F(z_2), z_1 - z_2 \rangle \geq \mu \|z_1 - z_2\|^2$$

Hence, if f is μ strongly-convex and strongly-concave and L -smooth, then the operator F is $2L$ -Lipschitz and μ strongly-monotone.

Gradient Descent Ascent for smooth, strongly-convex strongly-concave games

Claim: For L -smooth, μ strongly-convex strongly-concave games, T iterations of GDA with $\eta_k = \frac{\mu}{4L^2}$ results in the following bound,

$$\left\| \begin{bmatrix} w_T - w^* \\ v_T - v^* \end{bmatrix} \right\|^2 \leq \exp\left(\frac{-T}{4\kappa^2}\right) \left\| \begin{bmatrix} w_0 - w^* \\ v_0 - v^* \end{bmatrix} \right\|^2.$$

Proof: Since GDA can be equivalently written as $z_{k+1} = z_k - \eta F(z_k)$.

$$\begin{aligned} \|z_{k+1} - z^*\|^2 &= \|z_k - z^* - \eta F(z_k)\|^2 = \|z_k - z^*\|^2 - 2\eta \langle F(z_k), z_k - z^* \rangle + \eta^2 \|F(z_k)\|^2 \\ &= \|z_k - z^*\|^2 - 2\eta \langle F(z_k) - F(z^*), z_k - z^* \rangle + \eta^2 \|F(z_k) - F(z^*)\|^2 \\ &\quad (F(z^*) = 0 \text{ for unconstrained strongly-convex, strongly-concave games}) \\ &\leq \|z_k - z^*\|^2 - 2\eta \langle F(z_k) - F(z^*), z_k - z^* \rangle + 4L^2 \eta^2 \|z_k - z^*\|^2 \\ &\quad (F \text{ is } 2L\text{-Lipschitz}) \\ &\leq \|z_k - z^*\|^2 - 2\mu\eta \|z_k - z^*\|^2 + 4L^2 \eta^2 \|z_k - z^*\|^2 \\ &\quad (F \text{ is } \mu \text{ strongly-monotone}) \\ &= \|z_k - z^*\|^2 (1 - 2\mu\eta + 4L^2\eta^2) \end{aligned}$$

Gradient Descent Ascent for smooth, strongly-convex strongly-concave games

Recall that $\|z_{k+1} - z^*\|^2 \leq \|z_k - z^*\|^2 (1 - 2\mu\eta + 4L^2\eta^2)$. We need to set η such that $(1 - 2\mu\eta + 4L^2\eta^2) < 1 \implies \eta < \frac{\mu}{2L^2}$. Setting $\eta = \frac{\mu}{4L^2}$

$$\|z_{k+1} - z^*\|^2 \leq \|z_k - z^*\|^2 \left(1 - 2\mu \frac{\mu}{4L^2} + 4L^2 \frac{\mu^2 L^2}{16L^4}\right) \implies \|z_{k+1} - z^*\|^2 \leq \left(1 - \frac{\mu^2}{4L^2}\right) \|z_k - z^*\|^2$$

Recurring from $k = 0$ to $T - 1$,

$$\|z_T - z^*\|^2 \leq \left(1 - \frac{\mu^2}{4L^2}\right)^T \|z_0 - z^*\|^2 \leq \exp\left(\frac{-T}{4\kappa^2}\right) \|z_0 - z^*\|^2$$

(Since $1 - x \leq \exp(-x)$ for all x)

Hence, for smooth, strongly-convex strongly-concave games with condition number κ , we need to run GDA for $T = O(\kappa^2 \log(\frac{1}{\epsilon}))$ in order to get ϵ -close to the Nash equilibrium. The $O(\kappa^2)$ dependence can not be improved for GDA.

In contrast, for minimizing smooth, strongly-convex functions GD requires $O(\kappa \log(\frac{1}{\epsilon}))$ iterations in order to get ϵ -close to the minimizer.

Questions?