

Current Position

- **University of Alberta** Edmonton, Canada
Postdoctoral Researcher *November 2020 -*
 - Supervisor: Csaba Szepesvári

Education

- **University of British Columbia** Vancouver, Canada
Doctor of Philosophy (Computer Science) *Sep 2015 - Dec 2018*
 - Supervisors: Mark Schmidt, Laks Lakshmanan
 - Thesis: Structured Bandits and Applications
- **University of British Columbia** Vancouver, Canada
Master of Science (Computer Science) *Sep 2013 - July 2015*
 - Supervisor: Laks Lakshmanan
 - Thesis: Influence Maximization in Bandit and Adaptive settings
 - GPA: 4.32 / 4.33
- **Birla Institute of Technology and Science, Pilani** Goa, India
Bachelor of Engineering (Computer Science) *Aug 2008 - July 2012*
 - GPA: 9.37 / 10

Publications

- **Large-scale optimization**
 - [R] “SVRG meets AdaGrad: Painless Variance Reduction”, Benjamin Dubois-Taine*, **Sharan Vaswani***, Reza Babanezhad, Mark Schmidt, Simon Lacoste-Julien. arXiv, 2021. *Under conference submission.*
 - [W] “Adaptive Gradient Methods Converge Faster with Over-Parameterization (but you should do a line-search)”, **Sharan Vaswani**, Issam Laradji, Frederik Kunstner, Si Yi Meng, Mark Schmidt, Simon Lacoste-Julien. “Optimization for Machine Learning” workshop, NeurIPS, 2020 (**Spotlight**).
 - [C] “Stochastic Polyak Step-size for SGD: An Adaptive Learning Rate for Fast Convergence”, Nicolas Loizou, **Sharan Vaswani**, Issam Laradji, Simon Lacoste-Julien. International Conference on Artificial Intelligence and Statistics (AISTATS), 2021. “Optimization for Machine Learning” workshop, NeurIPS, 2020 (**Spotlight**).
 - [W] “How to make your optimizer generalize better”, **Sharan Vaswani**, Reza Babanezhad, Jose Gallego, Aaron Mishkin, Simon Lacoste-Julien, Nicolas Le Roux. “Optimization for Machine Learning” workshop, NeurIPS, 2020 (**Spotlight**).
 - [C] “Fast and Furious Convergence: Stochastic Second Order Methods under Interpolation”, Si Yi Meng*, **Sharan Vaswani***, Issam Laradji, Mark Schmidt, Simon Lacoste-Julien. International Conference on Artificial Intelligence and Statistics (AISTATS), 2020. “Beyond First Order Methods in Machine Learning” workshop, NeurIPS 2019 (**Spotlight**).
 - [C] “Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates”, **Sharan Vaswani**, Aaron Mishkin, Issam Laradji, Mark Schmidt, Gauthier Gidel, Simon Lacoste-Julien. Neural Information Processing Systems (NeurIPS), 2019.

C: Conference, W: Workshop, J: Journal, R: Technical report.

* Equal contribution.

- [W] “Accelerating boosting via accelerated greedy coordinate descent”, Xiaomeng Ju*, Yifan Sun*, **Sharan Vaswani***, Mark Schmidt. “Optimization for Machine Learning” workshop, NeurIPS 2019.
- [C] “Fast and Faster Convergence of SGD for Over-Parameterized Models and an Accelerated Perceptron”, **Sharan Vaswani**, Francis Bach, Mark Schmidt. International Conference on Artificial Intelligence and Statistics (AISTATS), 2019.

• Sequential decision-making

- [W] “A functional mirror ascent view of policy gradient methods with function approximation”, **Sharan Vaswani**, Olivier Bachem, Simone Totaro, Robert Müller, Matthieu Geist, Marlos Machado, Pablo Samuel Castro, Nicolas Le Roux, 2021. Workshop on Reinforcement Learning Theory, International Conference on Machine Learning (ICML), 2021. *Under conference submission.*
- [C] “Old Dog Learns New Tricks: Randomized UCB for Bandit Problems”, **Sharan Vaswani**, Abbas Mehrabian, Audrey Durand, Branislav Kveton. International Conference on Artificial Intelligence and Statistics (AISTATS), 2020.
- [J] “Combining Bayesian Optimization and Lipschitz Optimization”, Mohamed Osama Ahmed, **Sharan Vaswani**, Mark Schmidt. European Conference on Machine Learning (ECML) Journal Track, 2020.
- [C] “Garbage In, Reward Out: Bootstrapping Exploration in Multi-Armed Bandits”, Branislav Kveton, Csaba Szepesvári, **Sharan Vaswani**, Zheng Wen, Mohammad Ghavamzadeh, Tor Lattimore. International Conference on Machine Learning (ICML), 2019.
- [R] “New Insights into Bootstrapping for Bandits”, **Sharan Vaswani**, Branislav Kveton, Zheng Wen, Anup Rao, Mark Schmidt, Yasin Abbasi-Yadkori. arXiv:1805.09793, 2018.
- [C] “Online Influence Maximization under Independent Cascade Model with Semi-Bandit Feedback”, Zheng Wen, Branislav Kveton, Michal Valko, **Sharan Vaswani**. Neural Information Processing Systems (NIPS), 2017.
- [C] “Model-Independent Online Learning for Influence Maximization”, **Sharan Vaswani**, Branislav Kveton, Zheng Wen, Mohammad Ghavamzadeh, Laks Lakshmanan, Mark Schmidt. International Conference on Machine Learning (ICML), 2017.
- [C] “Horde of Bandits using Gaussian Markov Random Fields”, **Sharan Vaswani**, Mark Schmidt, Laks Lakshmanan. International Conference on Artificial Intelligence and Statistics (AISTATS), 2017. **(Oral presentation)**
- [W] “Influence Maximization with Bandits”, **Sharan Vaswani**, Laks Lakshmanan, Mark Schmidt. “Networks in Social and Information Sciences” workshop, NIPS, 2015.

• Social Networks

- [R] “Adaptive Influence Maximization in Social Networks: Why Commit when You can Adapt?”, **Sharan Vaswani**, Laks V.S. Lakshmanan. arXiv:1604.08171, 2016.
- [C] “Modeling Non-Progressive Phenomena for Influence Propagation”, Vincent Yun Lou, Smriti Bhagat, Laks Lakshmanan, **Sharan Vaswani**. ACM Conference on Online Social Networks (COSN), 2014.

• Parallel Computing

- [C] “Performance Evaluation of Medical Imaging Algorithms on Intel MIC Platform”, Jyotsna Khemka, Mrugesh Gajjar, **Sharan Vaswani**, Nagavijayalakshmi Vydyanathan, Rama Malladi, Vinutha V. IEEE International Conference on High Performance Computing (HiPC), 2013.
- [J] “Fast 3D Salient Region Detection in Medical Images using GPUs”, Thota, Rahul, **Sharan Vaswani**, Amit Kale, Nagavijayalakshmi Vydyanathan. Machine Intelligence and Signal Processing. Springer India, 2016.

- [C] “Fast 3D Structure Localization in Medical Volumes using CUDA-enabled GPUs”, **Sharan Vaswani**, Rahul Thota, Nagavijayalakshmi Vydyanathan, Amit Kale. IEEE International Conference on Parallel, Distributed and Grid Computing, 2012. (**Best paper award**)

Teaching & Supervision

- Teaching assistant:
 - Algorithms (2011), Theory of computation (2013), Computational optimization (2014), Artificial intelligence (2014).
 - Undergraduate Machine learning (2015, 2017, 2018).
 - Graduate Machine learning (2016, 2017).
- Students mentored/co-supervised:
 - Si Yi Meng, MSc, 2018-2020, University of British Columbia. Supervisor: Mark Schmidt.
Current position: PhD, Cornell University.
 - Aaron Mishkin, MSc, 2018-2020, University of British Columbia. Supervisor: Mark Schmidt.
Current position: PhD, Stanford University.
 - Benjamin Paul-Dubois-Taine, MSc, 2020-21, Paris-Saclay University. Supervisor: Alessandro Rudi.
Current position: PhD, Paris-Saclay University.
 - Jose Gallego, PhD, 2018-. Mila, Université de Montréal. Supervisor: Simon Lacoste-Julien.
 - Valentin Tiriach, MSc, 2020-, University of Alberta. Supervisor: Csaba Szepesvári.
 - Kenneth Tjhia, MSc, 2020-, University of Alberta. Supervisor: Csaba Szepesvári.
 - Shivam Garg, Research assistant, University of Alberta. Supervisor: Csaba Szepesvári.
 - Arushi Jain, PhD, 2018-, Mila, McGill University. Supervisor: Doina Precup.
 - Frederik Kunstner, PhD, 2019-, University of British Columbia. Supervisor: Mark Schmidt.
 - Jonathan Lavington, PhD, 2018-, University of British Columbia. Supervisor: Mark Schmidt.

Service

- Contributor to the Optimization chapter of the “Machine Learning, Second Edition: A Probabilistic Perspective” book by Kevin Murphy.
- Conference reviewer: AISTATS’19 - ’21, ICLR’18-’21, ICML’17-’20, JMLR’18-’21, IEEE TNNLS, NeurIPS’17-’21, New In ML workshop (NeurIPS’19), OPT-ML workshop (NeurIPS’20).
- Volunteer in the UBC Computer Science Graduate Admissions committee for 2016-2017, 2017-2018.
- Student representative in the UBC Computer Science Faculty Recruiting committee for 2015-2016.
- Conference volunteer for NIPS’16.
- Co-organizer of reading groups at UBC: Machine learning (2018), Deep learning (2015).
- Conference sub-reviewer for SIGMOD’18, AAAI’17,’18, WWW’17, SDM’15,’17, KDD’16,’17, ICDM’14.

Awards

- Postdoctoral Scholarship awarded by The Institute for Data Valorization (IVADO) (2019 - 2020).
- Four Year Doctoral Fellowship awarded by the University of British Columbia (2015 - 2018).

- Merit Scholarship awarded by the Birla Institute of Technology and Science, Pilani (2008-2010).
- Travel award for AISTATS 2017, ICML 2017-2019, NeurIPS 2017, 2019.
- Outstanding reviewer (Top 12%) for ICLR 2021, ICML 2021.
- Top 30% of highest scoring reviewers for NeurIPS 2018, 2019.

Employment

- **Mila, Université de Montréal** Montreal, Canada
Postdoctoral Researcher *January 2019 - October, 2020*
 - Supervisor: Simon Lacoste-Julien
- **Inria Paris** Paris, France
Intern *May 2018 - July, 2018*
 - Supervisor: Francis Bach
- **Apple** Seattle, USA
Intern *June 2017 - August, 2017*
 - Supervisors: Hoyt Koepke, Srikrishna Sridhar
- **Limespot** Vancouver, Canada
Machine Learning Consultant *March - May 2017; Sept, 2017 - Oct, 2017*
- **Adobe Research** San Jose, USA
Data Scientist Intern *Aug 2016 - Oct 2016*
 - Supervisors: Branislav Kveton, Zheng Wen, Mohammad Ghavamzadeh
- **University of British Columbia** Vancouver, Canada
Teaching Assistant *Sep 2013 - Dec 2018*
- **Siemens Corporate Research and Technologies** Bangalore, India
Research Engineer, Parallel Systems *July 2012 - June 2013*
 - Supervisors: Nagavijayalakshmi Vydyanathan, Amit Kale, Saptarshi Das
- **Siemens Corporate Research and Technologies** Bangalore, India
Research Intern, Parallel Systems *January 2012 - June 2012*
 - Supervisors: Nagavijayalakshmi Vydyanathan, Amit Kale
- **Birla Institute of Technology and Science, Pilani** Goa, India
Teaching Assistant *Aug 2012 - Dec 2012*
- **Indira Gandhi Centre for Atomic Research** Kalpakkam, India
Intern *May 2010 - July 2010*
 - Supervisor: M.L. Jayalal

Patents

- “Influence Maximization Determination in a Social Network System”, **Sharan Vaswani**, Branislav Kveton, Zheng Wen, Mohammad Ghavamzadeh. US Patent App. 15/611,597, 2018.

Talks

- “Rethinking Stochastic Optimization for Modern Machine Learning”
 - Simon Fraser University, Vancouver, March 2021.
 - York University, Toronto, March 2021.

- West Coast Optimization Meeting, Okanagan, May 2021.
- SIAM Conference on Optimization, July, 2021.
- “Old Dog Learns New Tricks: Randomized UCB for Bandit Problems”, Mila, Montreal, August 2020.
- “Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates”
 - Huawei Research, Montreal, October 2019.
 - Mila, Montreal, August 2019.
 - Element AI, Montreal, July 2019.
 - Google Brain, Montreal, July 2019.
- “Influence Maximization with Bandits”, UBC-Element AI workshop, August, 2018.
- “New Insights into Bootstrapping for Bandits”, Inria Paris, May 2018.
- “Model-Independent Online Learning for Influence Maximization”, International Conference on Machine Learning, 2017.
- “Horde of Bandits using Gaussian Markov Random Fields”, International Conference on Artificial Intelligence and Statistics, 2017.
- “Modeling Non-Progressive Phenomena for Influence Propagation” Conference on Online Social Networks, 2014.
- “Fast 3D Structure Localization in Medical Volumes using CUDA-enabled GPUs”, International Conference on Parallel, Distributed and Grid Computing, 2012.

References

- Mark Schmidt (schmidtm@cs.ubc.ca), Associate Professor, University of British Columbia.
- Csaba Szepesvári (csaba.szepesvari@gmail.com), Professor, University of Alberta.
- Nicolas Le Roux (nicolas.le.roux@gmail.com), Research Scientist, Microsoft Research.
- Simon Lacoste-Julien (slacoste@iro.umontreal.ca), Associate Professor, Université de Montréal.