

Sharan Vaswani
438, Seymour Street,
Vancouver, BC, Canada

Email: vaswani.sharan@gmail.com
Phone: +1 778-859-9246
Web: <https://vaswanis.github.io/>

Academic Employment

Simon Fraser University <i>Assistant Professor (Computing Science)</i>	Vancouver, Canada 2022 -
University of Alberta <i>Postdoctoral Researcher (Computing Science)</i> <ul style="list-style-type: none">◦ Supervisor: Csaba Szepesvári	Edmonton, Canada 2020 - 2021
Mila, Université de Montréal <i>Postdoctoral Researcher (Computer Science and Operations Research)</i> <ul style="list-style-type: none">◦ Supervisor: Simon Lacoste-Julien	Montreal, Canada 2019-2020

Education

University of British Columbia <i>Doctor of Philosophy (Computer Science)</i> <ul style="list-style-type: none">◦ Supervisors: Mark Schmidt, Laks Lakshmanan◦ Thesis: Structured Bandits and Applications	Vancouver, Canada 2015 - 2018
University of British Columbia <i>Master of Science (Computer Science)</i> <ul style="list-style-type: none">◦ Supervisor: Laks Lakshmanan◦ Thesis: Influence Maximization in Bandit and Adaptive settings◦ GPA: 4.32 / 4.33	Vancouver, Canada 2013 - 2015
Birla Institute of Technology and Science, Pilani <i>Bachelor of Engineering (Computer Science)</i> <ul style="list-style-type: none">◦ GPA: 9.37 / 10	Goa, India 2008 - 2012

Publications

Peer-reviewed Conference Publications

1. “Sample Complexity Bounds for Linear Constrained MDPs with a Generative Model”, Xingtú Liu, Lin F. Yang, **Sharan Vaswani**. *International Conference on Algorithmic Learning Theory (ALT)*, 2026.
“Constrained Optimization for Machine Learning” workshop, NeurIPS 2025.
2. “Preserving Plasticity in Continual Learning with Adaptive Linearity Injection”, Seyed Roozbeh Razavi Rohani, Khashayar Khajavi, Wesley Chung, Mo Chen, **Sharan Vaswani**. *Conference on Lifelong Learning Agents (CoLLAs)*, 2025.
3. “Armijo Line-search Can Make (Stochastic) Gradient Descent Provably Faster”, **Sharan Vaswani**, Reza Babanezhad. *International Conference on Machine Learning (ICML)*, 2025.
4. “Fast Convergence of Softmax Policy Mirror Ascent”, Reza Asad, Reza Babanezhad, Issam Laradji, Nicolas Le Roux, **Sharan Vaswani**. *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2025.
“Optimization for Machine Learning” workshop, NeurIPS, 2024.

Last updated: December 18, 2025

* Equal contribution.

5. “Small steps no more: Global convergence of stochastic gradient bandits for arbitrary learning rates”, Jincheng Mei, Bo Dai, Alekh Agarwal, **Sharan Vaswani**, Anant Raj, Csaba Szepesvari, Dale Schuurmans. *Neural Information Processing Systems (NeurIPS)*, 2024.
6. “Towards Principled, Practical Policy Gradient for Bandits and Tabular MDPs”, Michael Lu, Matin Aghaei, Anant Raj, **Sharan Vaswani**. *Reinforcement Learning Conference (RLC)*, 2024.
“Optimization for Machine Learning” workshop, NeurIPS, 2023 (**Oral Presentation**).
7. “From Inverse Optimization to Feasibility to ERM”, Saurabh Mishra, Anant Raj, **Sharan Vaswani**. *International Conference on Machine Learning (ICML)*, 2024.
“Optimization for Machine Learning” workshop, NeurIPS, 2023.
8. “Decision-Aware Actor-Critic with Function Approximation and Theoretical Guarantees”, **Sharan Vaswani**, Amirreza Kazemi, Reza Babanezhad, Nicolas Le Roux. *Neural Information Processing Systems (NeurIPS)*, 2023.
“Duality Principles for Modern ML” workshop, ICML, 2023.
9. “Target-based Surrogates for Stochastic Optimization”, Jonathan Lavington*, **Sharan Vaswani***, Reza Babanezhad, Mark Schmidt, Nicolas Le Roux. *International Conference on Machine Learning (ICML)*, 2023.
“Optimization for Machine Learning” workshop, NeurIPS, 2022.
10. “Near-Optimal Sample Complexity Bounds for Constrained MDPs”, **Sharan Vaswani***, Lin F. Yang*, Csaba Szepesvári. *Neural Information Processing Systems (NeurIPS)*, 2022.
11. “Towards Noise-adaptive, Problem-adaptive (Accelerated) Stochastic Gradient Descent”, **Sharan Vaswani**, Benjamin Dubois-Taine, Reza Babanezhad. *International Conference on Machine Learning (ICML)*, 2022 (**Oral Presentation**).
“Optimization for Machine Learning” workshop, NeurIPS, 2021.
12. “Improved Policy Optimization for Online Imitation Learning”, Jonathan Lavington, **Sharan Vaswani**, Mark Schmidt. *Conference on Lifelong Learning Agents (CoLLAs)*, 2022.
13. “Towards Painless Policy Optimization for Constrained MDPs”, Arushi Jain*, **Sharan Vaswani***, Reza Babanezhad, Doina Precup, Csaba Szepesvári, *Conference on Uncertainty in Artificial Intelligence (UAI)*, 2022.
14. “A general class of surrogate functions for stable and efficient reinforcement learning”, **Sharan Vaswani**, Olivier Bachem, Simone Totaro, Robert Müller, Shivam Garg, Matthieu Geist, Marlos Machado, Pablo Samuel Castro, Nicolas Le Roux. *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2022 (**Best Paper Honorable Mention**).
“Workshop on Reinforcement Learning Theory”, ICML 2021.
15. “Stochastic Polyak Step-size for SGD: An Adaptive Learning Rate for Fast Convergence”, Nicolas Loizou, **Sharan Vaswani**, Issam Laradji, Simon Lacoste-Julien. *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2021.
“Optimization for Machine Learning” workshop, NeurIPS, 2020 (**Spotlight**).
16. “Fast and Furious Convergence: Stochastic Second Order Methods under Interpolation”, Si Yi Meng*, **Sharan Vaswani***, Issam Laradji, Mark Schmidt, Simon Lacoste-Julien. *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2020.
“Beyond First Order Methods in Machine Learning” workshop, NeurIPS 2019 (**Spotlight**).
17. “Old Dog Learns New Tricks: Randomized UCB for Bandit Problems”, **Sharan Vaswani**, Abbas Mehrabian, Audrey Durand, Branislav Kveton. *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2020.
18. “Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates”, **Sharan Vaswani**, Aaron Mishkin, Issam Laradji, Mark Schmidt, Gauthier Gidel, Simon Lacoste-Julien. *Neural Information Processing Systems (NeurIPS)*, 2019.

19. "Garbage In, Reward Out: Bootstrapping Exploration in Multi-Armed Bandits", Branislav Kveton, Csaba Szepesvári, **Sharan Vaswani**, Zheng Wen, Mohammad Ghavamzadeh, Tor Lattimore. *International Conference on Machine Learning (ICML)*, 2019.
20. "Fast and Faster Convergence of SGD for Over-Parameterized Models and an Accelerated Perceptron", **Sharan Vaswani**, Francis Bach, Mark Schmidt. *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2019.
21. "Online Influence Maximization under Independent Cascade Model with Semi-Bandit Feedback", Zheng Wen, Branislav Kveton, Michal Valko, **Sharan Vaswani**. *Neural Information Processing Systems (NIPS)*, 2017.
22. "Model-Independent Online Learning for Influence Maximization", **Sharan Vaswani**, Branislav Kveton, Zheng Wen, Mohammad Ghavamzadeh, Laks Lakshmanan, Mark Schmidt. *International Conference on Machine Learning (ICML)*, 2017.
23. "Horde of Bandits using Gaussian Markov Random Fields", **Sharan Vaswani**, Mark Schmidt, Laks Lakshmanan. *International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2017. **(Oral Presentation)**
24. "Modeling Non-Progressive Phenomena for Influence Propagation", Vincent Yun Lou, Smriti Bhagat, Laks Lakshmanan, **Sharan Vaswani**. *ACM Conference on Online Social Networks (COSN)*, 2014.
25. "Performance Evaluation of Medical Imaging Algorithms on Intel MIC Platform", Jyotsna Khemka, Mrugesh Gajjar, **Sharan Vaswani**, Nagavijayalakshmi Vydyanathan, Rama Malladi, Vinutha V. *IEEE International Conference on High Performance Computing (HiPC)*, 2013.
26. "Fast 3D Structure Localization in Medical Volumes using CUDA-enabled GPUs", **Sharan Vaswani**, Rahul Thota, Nagavijayalakshmi Vydyanathan, Amit Kale. *IEEE International Conference on Parallel, Distributed and Grid Computing*, 2012. **(Best Paper Award)**

Peer-reviewed Journal Publications

1. "Noise-adaptive (Accelerated) Stochastic Heavy-Ball Momentum", Anh Dang, Reza Babanezhad, **Sharan Vaswani**, *Transactions on Machine Learning Research*, 2025.
"Optimization for Machine Learning" workshop, NeurIPS, 2023.
2. "SVRG meets AdaGrad: Painless Variance Reduction", Benjamin Dubois-Taine*, **Sharan Vaswani***, Reza Babanezhad, Mark Schmidt, Simon Lacoste-Julien. *Machine Learning Journal*, 2022.
3. "Combining Bayesian Optimization and Lipschitz Optimization", Mohamed Osama Ahmed, **Sharan Vaswani**, Mark Schmidt. *Machine Learning Journal*, 2020.
4. "Fast 3D Salient Region Detection in Medical Images using GPUs", Thota, Rahul, **Sharan Vaswani**, Amit Kale, Nagavijayalakshmi Vydyanathan. *Machine Intelligence and Signal Processing. Springer India*, 2016.

Peer-reviewed Workshop Publications

1. "Optimistic Actor-Critic with Parametric Policies: Unifying Sample Efficiency and Practicality", Max Qiushi Lin, Reza Asad, Kevin Tan, Haque Ishfaq, Csaba Szepesvari, **Sharan Vaswani**, Workshop on Aligning Reinforcement Learning Experimentalists and Theorists (ARLET), NeurIPS, 2025. *(Under conference submission)*
2. "Towards Parameter-Free Temporal Difference Learning", Yunxiang Li, Mark Schmidt, Reza Babanezhad, **Sharan Vaswani**, Workshop on Aligning Reinforcement Learning Experimentalists and Theorists (ARLET), NeurIPS, 2025. *(Under conference submission)*
3. "Revisiting Actor-Critic Methods in Discrete Action Off-Policy Reinforcement Learning", Reza Asad, Reza Babanezhad, **Sharan Vaswani**, Workshop on Aligning Reinforcement Learning Experimentalists and Theorists (ARLET), NeurIPS, 2025. *(Under conference submission)*

4. “Implicit Bias of Polyak and Line-Search Step Sizes on Linear Classification with Separable Data”, Chen Fan, Reza Babanezhad, Christos Thrampoulidis, Mark Schmidt, **Sharan Vaswani**, Optimization for Machine Learning (OPT) workshop, NeurIPS, 2025.
5. “Improving OOD Generalization of Pre-trained Encoders via Aligned Embedding-Space Ensembles”, Shuman Peng, Arash Khoeini, **Sharan Vaswani**, Martin Ester, “Unifying Representations in Neural Models (UniReps)” workshop, “Self-supervised Learning (SSL) workshop”, NeurIPS, 2024.
6. “Surrogate Minimization: An Optimization Algorithm for Training Large Neural Networks with Model Parallelism”, Reza Asad, Reza Babanezhad, Issam Laradji, Nicolas Le Roux, **Sharan Vaswani**. “Optimization for Machine Learning” workshop, NeurIPS, 2023.
7. “MSL: An Adaptive Momentum-based Stochastic Line-search Framework”, Chen Fan, **Sharan Vaswani**, Christos Thrampoulidis, Mark Schmidt. “Optimization for Machine Learning” workshop, NeurIPS, 2023.
8. “Adaptive Gradient Methods Converge Faster with Over-Parameterization (but you should do a line-search)”, **Sharan Vaswani**, Issam Laradji, Frederik Kunstner, Si Yi Meng, Mark Schmidt, Simon Lacoste-Julien. “Optimization for Machine Learning” workshop, NeurIPS, 2020 (**Spotlight**).
9. “How to Make Your Optimizer Generalize Better”, **Sharan Vaswani**, Reza Babanezhad, Jose Gallego, Aaron Mishkin, Simon Lacoste-Julien, Nicolas Le Roux. “Optimization for Machine Learning” workshop, NeurIPS, 2020 (**Spotlight**).
10. “Accelerating boosting via accelerated greedy coordinate descent”, Xiaomeng Ju*, Yifan Sun*, **Sharan Vaswani***, Mark Schmidt. “Optimization for Machine Learning” workshop, NeurIPS 2019.
11. “Influence Maximization with Bandits”, **Sharan Vaswani**, Laks Lakshmanan, Mark Schmidt. “Networks in Social and Information Sciences” workshop, NIPS, 2015.

Preprints

1. “Glocal Smoothness: Line Search can really help!”, Curtis Fox, Aaron Mishkin, **Sharan Vaswani**, Mark Schmidt. ArXiv, 2025. (*Under journal submission*).
2. “Rethinking the Global Convergence of Softmax Policy Gradient with Linear Function Approximation”, Max Qiushi Lin, Jincheng Mei, Matin Aghaei, Michael Lu, Bo Dai, Alekh Agarwal, Dale Schuurmans, Csaba Szepesvari, **Sharan Vaswani**. ArXiv, 2025. (*Under journal submission*).
3. “New Insights into Bootstrapping for Bandits”, **Sharan Vaswani**, Branislav Kveton, Zheng Wen, Anup Rao, Mark Schmidt, Yasin Abbasi-Yadkori. ArXiv, 2018.
4. “Adaptive Influence Maximization in Social Networks: Why Commit when You can Adapt?”, **Sharan Vaswani**, Laks V.S. Lakshmanan. ArXiv, 2016.

Awards

- Discovery Grant awarded by the Natural Sciences and Engineering Research Council of Canada (2022-2027) [157,500 CAD for 5 years].
- Postdoctoral Scholarship awarded by The Institute for Data Valorization (IVADO) (2019 - 2020).
- Four Year Doctoral Fellowship awarded by the University of British Columbia (2015 - 2018).
- Merit Scholarship awarded by the Birla Institute of Technology and Science, Pilani (2008-2010).
- Travel award for AISTATS 2017, ICML 2017-2019, NeurIPS 2017, 2019.
- Outstanding reviewer for ICLR 2021, ICML 2021, 2022.
- Top 30% of highest scoring reviewers for NeurIPS 2018, 2019.

Teaching

Simon Fraser University

Instructor

2022-

- CMPT 210: Probability and Computing (undergraduate course)
 - * Summer, 2022. Class size: 50.
 - * Spring, 2023. Class size: 41.
 - * Spring, 2024. Class size: 71.
 - * Fall, 2024. Class size: 199.
 - * Fall, 2025. Class size: 200.
- CMPT 409/981: Optimization for Machine Learning (cross-listed undergraduate/graduate course)
 - * Fall, 2022. Class size: 33 (5 undergraduate students, 28 graduate students).
 - * Fall, 2024. Class size: 23 (9 undergraduate students, 14 graduate students).
 - * Fall, 2025. Class size: 18 (8 undergraduate students, 10 graduate students).
- CMPT 419/983: Theoretical Foundations of Reinforcement Learning (cross-listed undergraduate/graduate course)
 - * Fall, 2023. Class size: 25 (6 undergraduate students, 19 graduate students).

University of British Columbia

Teaching Assistant

2013-2018

- Theory of Computation (2013), Computational Optimization (2014), Artificial Intelligence (2014).
- Undergraduate Machine learning (2015, 2017, 2018).
- Graduate Machine learning (2016, 2017).

Birla Institute of Technology and Science, Pilani

Teaching Assistant

2011

- Data Structures and Algorithms (2011)

Supervision

- Students supervised (as the primary supervisor at Simon Fraser University)
 - Qiushi Lin, PhD
 - Ahmed Shehata, PhD
 - Reza Asad, PhD (co-supervised with Manolis Savva)
 - Michael Lu, PhD (co-supervised with Mo Chen)
 - Xingtu Liu, MSc, 2025. *Thesis:* Sample Complexity Bounds for Linear Constrained MDPs with a Generative Model.
 - Matin Aghaei, MSc, 2025. *Thesis:* Convergence of Softmax Policy Gradient: Incorporating Entropy Regularization and Handling Linear Function Approximation. *Next position:* Huawei
 - Anh Dang, MSc, 2024. *Thesis:* (Accelerated) Noise-adaptive Stochastic Heavy-Ball Momentum. *Next position:* Amazon
 - Saurabh Mishra, MSc, 2024. *Thesis:* From Inverse Optimization to Feasibility to Empirical Risk Minimization. *Next position:* Huawei

- Students mentored/unofficially co-supervised (as Assistant Professor):
 - Yunxiang Li, PhD, University of British Columbia. Supervisor: Mark Schmidt.
 - Chen Fan, PhD, University of British Columbia. Supervisors: Mark Schmidt, Christos Thrampoulidis.
 - Amirreza Kazemi, MSc, Simon Fraser University. Supervisor: Martin Ester.
Next position: Huawei
 - Jonathan Lavington, PhD, University of British Columbia. Supervisor: Mark Schmidt.
Next position: Amazon.
 - Valentin Tiriac, MSc, University of Alberta. Supervisor: Csaba Szepesvári.
 - Arushi Jain, PhD, Mila, McGill University. Supervisor: Doina Precup.
- Students mentored/unofficially co-supervised (as Postdoctoral Researcher):
 - Si Yi Meng, MSc, University of British Columbia. Supervisor: Mark Schmidt.
Next position: PhD, Cornell University.
 - Aaron Mishkin, MSc, University of British Columbia. Supervisor: Mark Schmidt.
Next position: PhD, Stanford University.
 - Benjamin Paul-Dubois-Taine, MSc, Paris-Saclay University. Supervisor: Alessandro Rudi.
Next position: PhD, Paris-Saclay University.
 - Shivam Garg, Research assistant, University of Alberta. Supervisor: Csaba Szepesvári.
Next position: PhD, University of Alberta.
 - Jose Gallego, PhD, Mila, Université de Montréal. Supervisor: Simon Lacoste-Julien.
 - Frederik Kunstner, PhD, University of British Columbia. Supervisor: Mark Schmidt.
Next position: Postdoc, Inria Paris.
- Committee Member (at Simon Fraser University)
 - Shuman Peng, Ph.D. Supervisor: Martin Ester.
 - Sriraj Meenavilli, MSc, 2023. Supervisor: Mo Chen.
 - Amirreza Kazemi, MSc, 2023. Supervisor: Martin Ester
- Examiner (at Simon Fraser University)
 - Seyed Roozbeh Razavi Rohani, PhD. Supervisor: Mo Chen
 - Dekai Lin, MSc, 2023. Supervisor: Ke Wang
- External Examiner
 - Stephen Scinocca, MSc, 2022. University of Victoria. Supervisor: Nishant Mehta.

Internal Service

- Member of Undergraduate Program committee (Jan'22 - August'22), Graduate Program committee (September'22 - Present) at Simon Fraser University.
- Volunteer in the UBC Computer Science Graduate Admissions committee for 2016-2017, 2017-2018.
- Student representative in the UBC Computer Science Faculty Recruiting committee for 2015-2016.
- Co-organizer of reading groups at UBC: Machine learning (2018), Deep learning (2015).

External Service

- Organizer of the “Mathematics of Machine Learning” research theme at the 2022 annual meeting of the Canadian Applied and Industrial Mathematics Society.
- External Reviewer for NSERC Arthur B. McDonald Fellowship, 2025.
- Conference Area Chair: NeurIPS’22-’25; ICML’23-’26; AAAI’25-’26, RLC’25
- Reviewer: JMLR’18-’23, AISTATS’19 - ’21, ICLR’18-’21, ICML’17-’22, IEEE TNNLS, NeurIPS’17-’21, New In ML workshop (NeurIPS’19), OPT-ML workshop (NeurIPS’20 - ’21).
- Conference volunteer for NIPS’16.
- Conference sub-reviewer for SIGMOD’18, AAAI’17,’18, WWW’17, SDM’15,’17, KDD’16,’17, ICDM’14.
- Contributor to the Optimization chapter of the “Machine Learning, Second Edition: A Probabilistic Perspective” book by Kevin Murphy.

Other Employment

Inria Paris

Intern

Paris, France

May - July, 2018

- Supervisor: Francis Bach

Apple

Intern

Seattle, USA

June - August, 2017

- Supervisors: Hoyt Koepke, Srikrishna Sridhar

Limespot

Machine Learning Consultant

Vancouver, Canada

March - May 2017; Sept - Oct, 2017

Adobe Research

Data Scientist Intern

San Jose, USA

August - Oct, 2016

- Supervisors: Branislav Kveton, Zheng Wen, Mohammad Ghavamzadeh

Siemens Corporate Research and Technologies

Research Engineer, Parallel Systems

Bangalore, India

July 2012 - June 2013

- Supervisors: Nagavijayalakshmi Vydyanathan, Amit Kale, Saptarshi Das

Research Intern, Parallel Systems

January - June, 2012

- Supervisors: Nagavijayalakshmi Vydyanathan, Amit Kale

Indira Gandhi Centre for Atomic Research

Intern

Kalpakkam, India

May 2010 - July 2010

- Supervisor: M.L. Jayalal

Patents

- “Influence Maximization Determination in a Social Network System”, **Sharan Vaswani**, Branislav Kveton, Zheng Wen, Mohammad Ghavamzadeh. US Patent App. 15/611,597, 2018.

Invited Talks

- “A Systematic Framework for Designing Policy Gradient Methods: The Case of Softmax Policy Mirror Ascent”
 - Google DeepMind, London, June 2025.
- “From Inverse Optimization to Feasibility to ERM”
 - INFORMS Computing Society Conference (ICS), Toronto, March 2025.
 - INFORMS Annual Meeting, Seattle, October 2024.
- “Global Convergence of Softmax Policy Gradient for Stochastic Bandits”
 - Canadian Mathematical Society Winter Meeting, Vancouver, November 2024.
- “Towards Principled, Practical Policy Gradient for Bandits and Tabular MDPs”
 - Theory of RL Workshop, Alberta, June 2024.
- “A Walk with SGD: Interpolation, Problem and Noise Adaptivity”
 - Department of Combinatorics and Optimization, University of Waterloo, January, 2024.
- “Exploiting Problem Structure for Efficient Optimization in Machine Learning”
 - Operations Research Seminar, Simon Fraser University, October, 2023.
- “Decision-Aware Actor-Critic with Function Approximation and Theoretical Guarantees”
 - Vector Institute, Toronto, June, 2023.
- “Target-based Surrogates for Efficient Sequential decision-making”
 - Microsoft Research, Montreal, June, 2023.
 - AI Seminar, Simon Fraser University, March 2023.
- “A general class of surrogate functions for stable and efficient reinforcement learning”
 - Theory of RL Workshop, Alberta, May 2023.
 - RL Seminar, Mila, Montreal, March 2022.
 - Facebook AI Research, Paris, April 2022.
- “Towards Noise-adaptive, Problem-adaptive Stochastic Gradient Descent”
 - Siam Conference on Optimization, Seattle, June 2023.
 - Optimization Seminar, Mila, Montreal, November 2021.
 - RWTH Aachen University, April, 2022.
- “Rethinking Stochastic Optimization for Modern Machine Learning”
 - Simon Fraser University, Vancouver, March 2021.

- York University, Toronto, March 2021.
- West Coast Optimization Meeting, Okanagan, May 2021.
- SIAM Conference on Optimization, July, 2021.
- “Old Dog Learns New Tricks: Randomized UCB for Bandit Problems”, RL Seminar, Mila, Montreal, August 2020.
- “Painless Stochastic Gradient: Interpolation, Line-Search, and Convergence Rates”
 - Huawei Research, Montreal, October 2019.
 - Optimization Seminar, Mila, Montreal, August 2019.
 - Element AI, Montreal, July 2019.
 - Google Brain, Montreal, July 2019.
- “Influence Maximization with Bandits”, UBC-Element AI workshop, August, 2018.
- “New Insights into Bootstrapping for Bandits”, Inria Paris, May 2018.